

AntiCopy

I built **AntiCopy**, a small helper tool for interviewers. It sits quietly next to a video call (like Google Meet) and watches how a candidate behaves. It does not join the meeting, does not install anything on the candidate's machine, and it never says "this person is cheating." It only shows a risk score, some evidence, and a few suggestions. The interviewer always decides what to do with that.

Things Used:

On the browser side I used **Next.js** for the website and dashboard. For the candidate's face on the Meet tab I used **MediaPipe Face Landmarker**- it tracks eye gaze, blinks, and whether they keep looking down or off-screen (like reading from another screen). For speech I used the browser's **Web Speech API** so I could get a live transcript and notice things like how fast they talk, how long they take to answer, pauses, and filler words. The interviewer also shares the meeting tab with audio so I can hear the candidate.

On the server side I used **FastAPI (Python)**. The browser sends small bits of data over a **WebSocket**. Then my feature extractors pull out timing, speech rate, language patterns, ownership of answers, and gaze stats. Separate detectors look for two kinds of signals: ones that raise suspicion (like very flat answer times), and ones that look human (like natural variation). A risk engine subtracts the human signals from the suspicious ones, so one weird glance cannot decide the whole score. A flow engine suggests what kind of question to ask next.

Idea:

I thought about it as a candidate. Not that I've used AI tools, but how'd I convince the interviewer FOR SURE I'm not using AI?

If someone answers every question after almost the same delay, speaks at a very steady speed, gives generic answers, or keeps looking down like they are reading, the score can go up. If they hesitate, change pace, tell a real story from their own work, ask for clarity, or sound like they are thinking out loud, those human signals pull the score back down (for NLP `sPaCy` was used). That push-and-pull is intentional — I would rather miss something than wrongly flag an honest candidate.

Features:

Eye gaze is used as supporting evidence only. It's only useful when it lines up with other signs, never enough on its own since ONLY eye gaze is something tough to infer with that candidate's cheating. Voice modulation (how speech speed and rhythm change) helps tell reading apart from real thinking. Question suggestion tells the interviewer what to try next based on what the session is showing. The specialty of rapid easy questions is simple: fire a quick factual question and watch the delay- a live AI round-trip often stays oddly flat even on easy stuff. Opinion-based questions matter because there is no neat lookup answer; you have to reason in your own words. Cross-questioning and confirmation (follow-ups that check or clash with what they just said) put pressure on a scripted reply. That clash of voices- clarification, pushback, personal ownership versus smooth generic answers- is the human side of the system: real people leave friction; assistants often sound too even.

Coverage and limits (honest note)

AntiCopy covers browser-only monitoring for remote interviews when the Meet tab (with audio) is shared and the candidate's face is visible. It is stronger when several signals agree over a longer chat, and weaker at the start when there is little data. Transcription quality depends on the browser and the meeting audio. Gaze needs a clear face on screen. It will not catch every AI tool, and it can be wrong. It estimates probability from behaviour. It is a reviewer aid, not a judge, and it should never be used as sole proof against a candidate. The novel ideas like quick question can have issues like if interviewee actually lags or network issues. It also needs at least 10 minutes of interview for score to get high, for 5-6 minutes, the score shows low AI Risk scores but one can STILL evaluate using reasons why it is and it isn't AI.